


# From APEX.AMP To torch.cuda.AMP

Hyungon Ryu | NVAITC Korea



# AGENDA


**Precision**

**apex.amp**


**torch.cuda.amp**

**PRECISION**

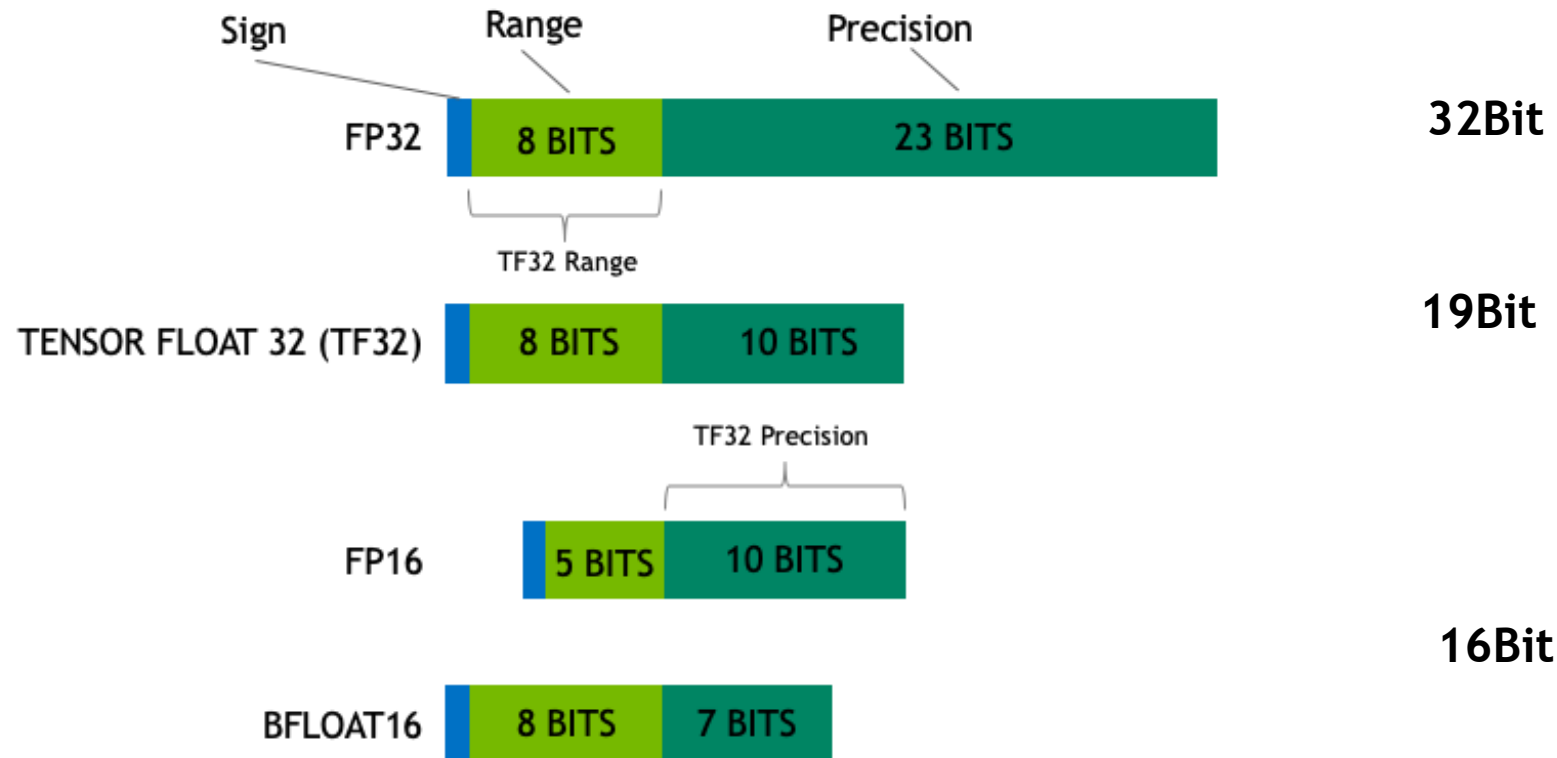
# FP32 TRAINING



# FP16 TRAINING




# PRECISION FOR DL



**Benefit of FP16**  
Lowering required memory  
~ large mini-batches  
~ less data transfer  
Faster training/inferencing


# A HISTOGRAM OF ACTIVATION GRADIENTS

Representation of Half Precision floating point format



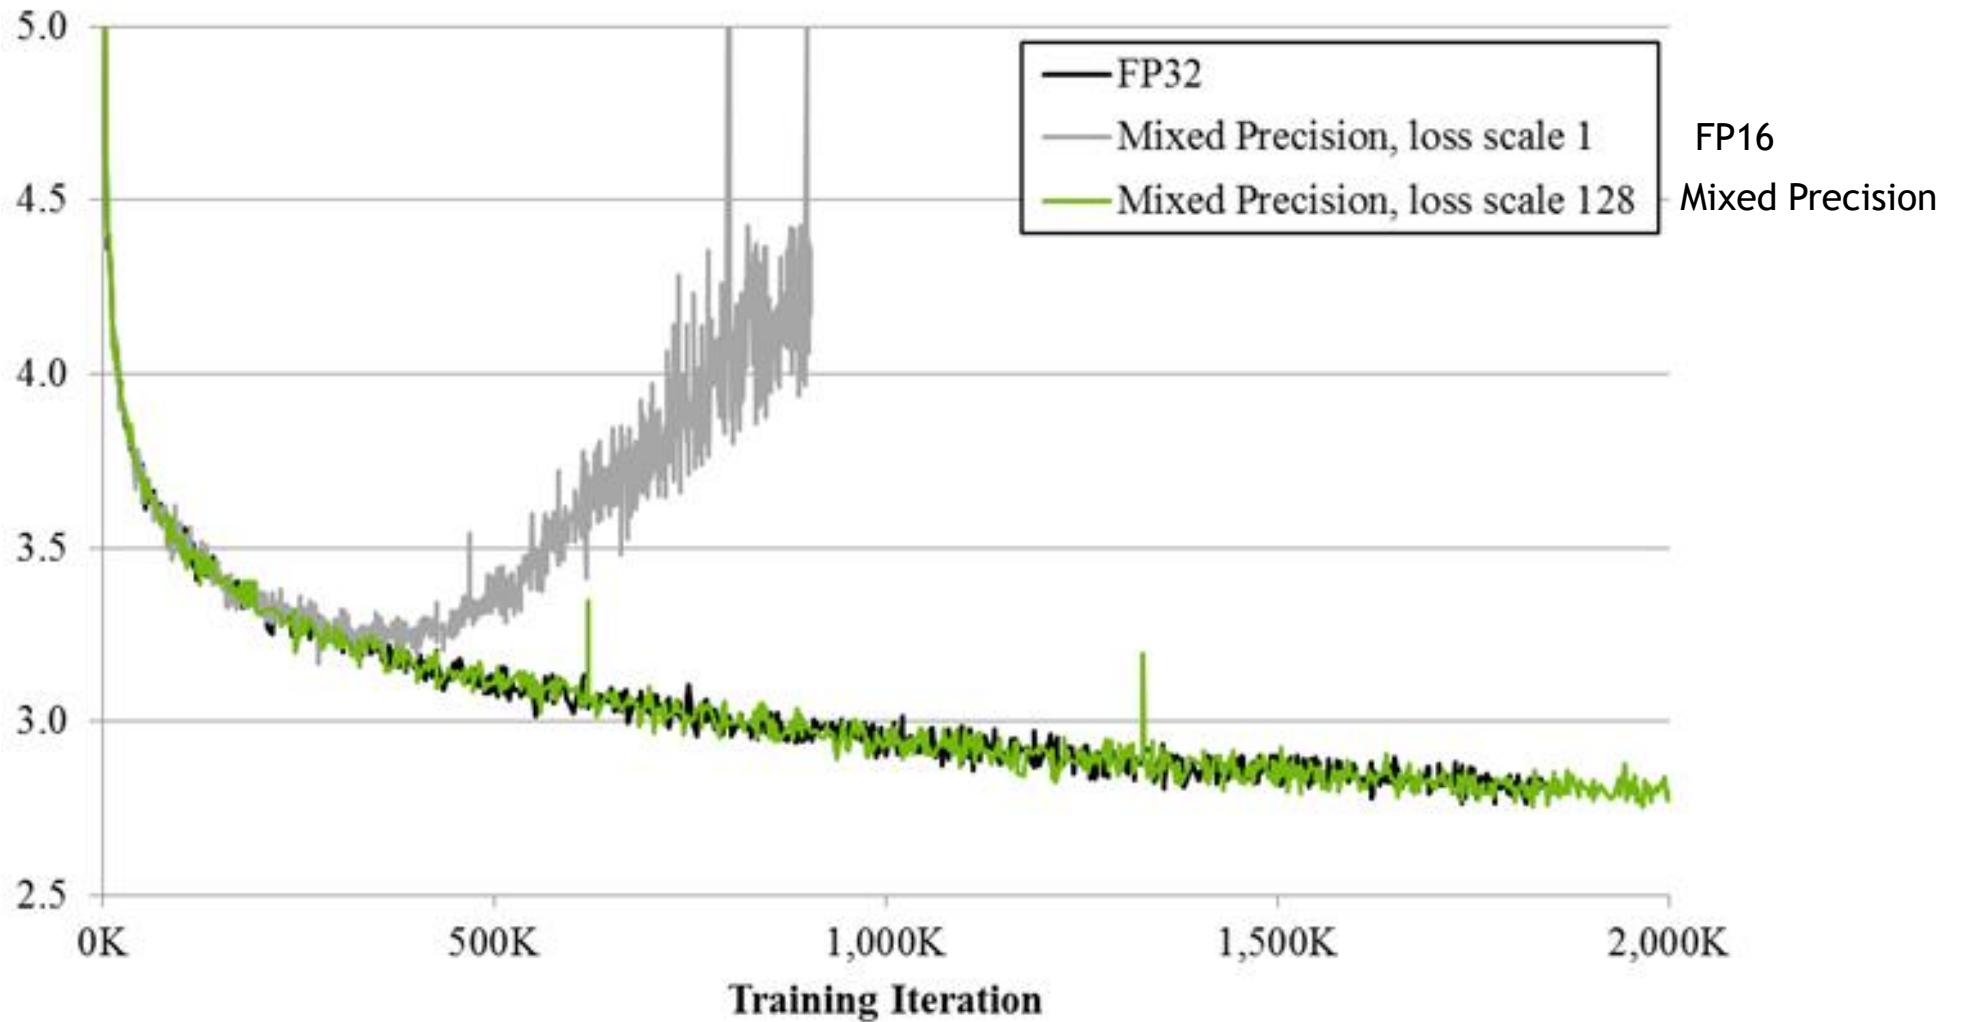
FP32		FP16
0.0000001	→	0
Near Zero		Zero

# GRADIENTS MAY UNDERFLOW






# GRADIENT EXPLODE



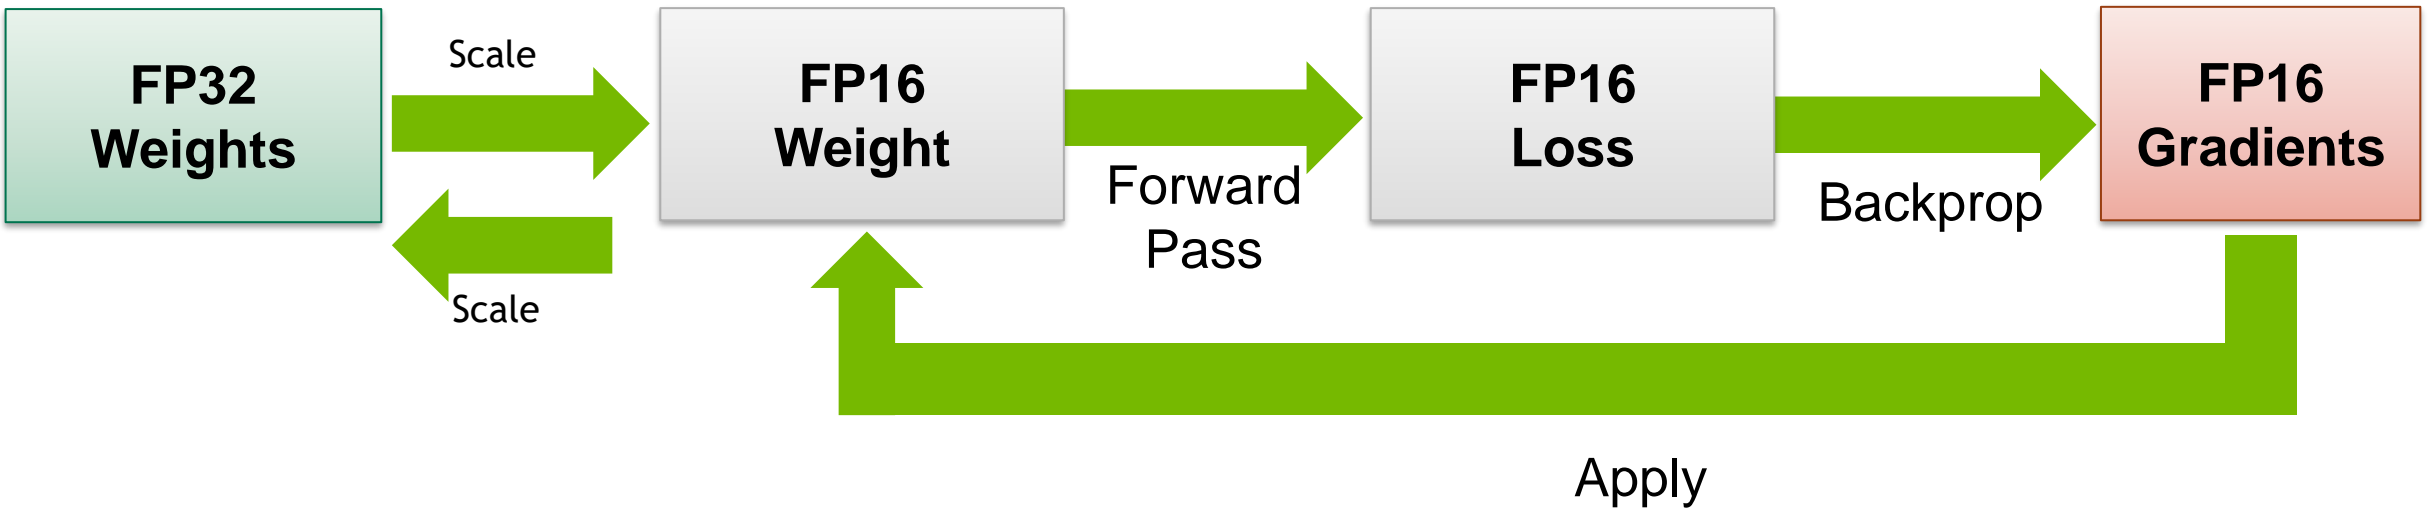
# MIXED PRECISION TRAINING

Compute in FP16  
Save in FP32



# MIXED PRECISION TRAINING

Compute in FP16  
Save in FP32



cast to fp16  
 $\times \text{loss\_scale}$

cast to fp32  
 $\div \text{loss\_scale}$

**APEX.AMP**

# APEX.AMP

<https://nvidia.github.io/apex/amp.html>

APEX enable DDP, Mixed Precision(FP16) and AMP

```
amp.initialize(model, optimizer, opt_level="O1")
```

O0 : FP32, O1/O2 : AMP, O3 : FP16

```
with amp.scale_loss(loss, optimizer) as scaled_loss:
```

```
    scaled_loss.backward()
```

# **3 STEP FOR APEX.AMP**

**Load APEX.AMP module**

**Initiate APEX AMP**

**Enable Loss scale**

# APEX.AMP

## Step1. load APEX module

```
39 import torch.distributed as dist
40 from torch.utils.data.distributed import DistributedSampler
41
42 from apex.parallel import DistributedDataParallel as DDP
43
44 import models
45 import loss_functions
46 import data_functions
47 from tacotron2_common.utils import ParseFromConfigFile
48
49 import dllogger as DLLogger
50 from dllogger import StdOutBackend, JSONStreamBackend, Verbosity
51
52 from scipy.io.wavfile import write as write_wav
53
54 from apex import amp
55 amp.lists.functional_overrides.FP32_FUNCS.remove('softmax')
56 amp.lists.functional_overrides.FP16_FUNCS.append('softmax')
57
```



# APEX.AMP

## Step2. initiate AMP

```
387 model_config = models.get_model_config(model_name, args)
388 model = models.get_model(model_name, model_config,
389                          cpu_run=False,
390                          uniform_initialize_bn_weight=not args.disable_uniform_init:
391
392 if not args.amp and distributed_run:
393     model = DDP(model)
394
395 optimizer = torch.optim.Adam(model.parameters(), lr=args.learning_rate,
396                               weight_decay=args.weight_decay)
397
398 if args.amp:
399     model, optimizer = amp.initialize(model, optimizer, opt_level="O1")
400     if distributed_run:
401         model = DDP(model)
402
```






# APEX.AMP

## Step3. enable loss scale

```
481 x, y, num_items = batch_to_gpu(batch)
482
483 y_pred = model(x)
484 loss = criterion(y_pred, y)
485
486 if distributed_run:
487     reduced_loss = reduce_tensor(loss.data, world_size).item()
488     reduced_num_items = reduce_tensor(num_items.data, 1).item()
489 else:
490     reduced_loss = loss.item()
491     reduced_num_items = num_items.item()
```



```
502 if args.amp:
503     with amp.scale_loss(loss, optimizer) as scaled_loss:
504         scaled_loss.backward()
505         grad_norm = torch.nn.utils.clip_grad_norm_(
506             amp.master_params(optimizer), args.grad_clip_thresh)
507 else:
508     loss.backward()
509     grad_norm = torch.nn.utils.clip_grad_norm_(
510         model.parameters(), args.grad_clip_thresh)
511
512 optimizer.step()
```

# ALSO NEED

## amp.state\_dict()

```
191 def save_checkpoint(model, optimizer, epoch, config, amp_run, output_dir, model_name,
192                     local_rank, world_size):
193
194     random_rng_state = torch.random.get_rng_state().cuda()
195     cuda_rng_state = torch.cuda.get_rng_state(local_rank).cuda()
196
197     random_rng_states_all = [torch.empty_like(random_rng_state) for _ in range(world_size)]
198     cuda_rng_states_all = [torch.empty_like(cuda_rng_state) for _ in range(world_size)]
199
200     if world_size > 1:
201         dist.all_gather(random_rng_states_all, random_rng_state)
202         dist.all_gather(cuda_rng_states_all, cuda_rng_state)
203     else:
204         random_rng_states_all = [random_rng_state]
205         cuda_rng_states_all = [cuda_rng_state]
206
207     random_rng_states_all = torch.stack(random_rng_states_all).cpu()
208     cuda_rng_states_all = torch.stack(cuda_rng_states_all).cpu()
209
210     if local_rank == 0:
211         checkpoint = {'epoch': epoch,
212                      'cuda_rng_state_all': cuda_rng_states_all,
213                      'random_rng_states_all': random_rng_states_all,
214                      'config': config,
215                      'state_dict': model.state_dict(),
216                      'optimizer': optimizer.state_dict()}
217
218         if amp_run:
219             checkpoint['amp'] = amp.state_dict()
```

```
246 def load_checkpoint(model, optimizer, epoch, config, amp_run, filepath, local_rank):
247
248     checkpoint = torch.load(filepath, map_location='cpu')
249
250     epoch[0] = checkpoint['epoch']+1
251     device_id = local_rank % torch.cuda.device_count()
252     torch.cuda.set_rng_state(checkpoint['cuda_rng_state_all'][device_id])
253     if 'random_rng_states_all' in checkpoint:
254         torch.random.set_rng_state(checkpoint['random_rng_states_all'][device_id])
255     elif 'random_rng_state' in checkpoint:
256         torch.random.set_rng_state(checkpoint['random_rng_state'])
257     else:
258         raise Exception("Model checkpoint must have either 'random_rng_state' or 'random_rng_states_all' key.")
259     config = checkpoint['config']
260     model.load_state_dict(checkpoint['state_dict'])
261     optimizer.load_state_dict(checkpoint['optimizer'])
262
263     if amp_run:
264         amp.load_state_dict(checkpoint['amp'])
265
```

**TORCH.CUDA.AMP**

# TORCH NATIVE AMP

<https://pytorch.org/docs/stable/amp.html>

APEX.AMP Upstream to pytorch ( Oct. 2020)

```
torch.cuda.amp.autocast(enabled=True)
```

```
torch.cuda.amp.GradScaler(init_scale=65536.0, growth_factor=2.0, backoff_factor=0.5, growth_interval=2000, enabled=True)
```

# TORCH.CUDA.AMP

## Step1. enable scaler

[illegible]

# TORCH.CUDA.AMP

## Step2. enable autocast



```
464 def train_iteration(model, i, mems, data_chunks, target_chunks, scaler,
465                     optimizer, device, delay_unscale, args):
466     cpu = torch.device('cpu')
467     data_i = data_chunks[i].contiguous()
468     target_i = target_chunks[i].contiguous()
469
470     if args.swap_mem and mems[i] is not None:
471         mems[i] = mems[i].to(device, non_blocking=True)
472
473     enable_autocast = args.fp16 and args.amp == 'pytorch'
474     with torch.cuda.amp.autocast(enable_autocast):
475         loss, mems[i] = model(data_i, target_i, mems[i])
476         loss = loss.float().mean().type_as(loss) / args.batch_chunk
477
478     if args.swap_mem and mems[i] is not None:
479         mems[i] = mems[i].to(cpu, non_blocking=True)
480
481     if args.fp16:
482         if args.amp == 'pytorch':
483             scaler.scale(loss).backward()
484         elif args.amp == 'apex':
485             with amp.scale_loss(loss, optimizer, delay_unscale=delay_unscale) as scaled_loss:
486                 scaled_loss.backward()
487     else:
488         loss.backward()
489
490     train_loss = loss.float().item()
491     return train_loss
492
```

# TORCH.CUDA.AMP

## Step3. use scaler

```
464 def train_iteration(model, i, mems, data_chunks, target_chunks, scaler,
465                     optimizer, device, delay_unscale, args):
466     cpu = torch.device('cpu')
467     data_i = data_chunks[i].contiguous()
468     target_i = target_chunks[i].contiguous()
469
470     if args.swap_mem and mems[i] is not None:
471         mems[i] = mems[i].to(device, non_blocking=True)
472
473     enable_autocast = args.fp16 and args.amp == 'pytorch'
474     with torch.cuda.amp.autocast(enable_autocast):
475         loss, mems[i] = model(data_i, target_i, mems[i])
476         loss = loss.float().mean().type_as(loss) / args.batch_chunk
477
478     if args.swap_mem and mems[i] is not None:
479         mems[i] = mems[i].to(cpu, non_blocking=True)
480
481     if args.fp16:
482         if args.amp == 'pytorch':
483             scaler.scale(loss).backward()
484         elif args.amp == 'apex':
485             with amp.scale_loss(loss, optimizer, delay_unscale=delay_unscale) as scaled_loss:
486                 scaled_loss.backward()
487     else:
488         loss.backward()
489
490     train_loss = loss.float().item()
491     return train_loss
492
```



# ALSO NEED

## scaler.state\_dict()



```
---
296 def save_checkpoint(args, model, model_config, optimizer, scheduler, scaler,
297                     vocab, epoch, batch, last_iter, train_step, best_val_loss,
298                     is_best, work_dir):
299     if args.fp16:
300         if args.amp == 'pytorch':
301             amp_state = scaler.state_dict()
302         elif args.amp == 'apex':
303             amp_state = amp.state_dict()
304     else:
305         amp_state = None
306
307     state = {
308         'args': args,
309         'model_config': model_config,
310         'model_state': model.state_dict(),
311         'optimizer_state': optimizer.state_dict(),
312         'scheduler_state': scheduler.state_dict(),
313         'vocab': vocab,
314         'amp_state': amp_state,
315         'epoch': epoch,
316         'batch': batch,
317         'last_iter': last_iter,
318         'train_step': train_step,
319         'best_val_loss': best_val_loss,
320     }
---
```



# ADDITIONAL TIPS

Do not check `loss.item()` for every iteration

