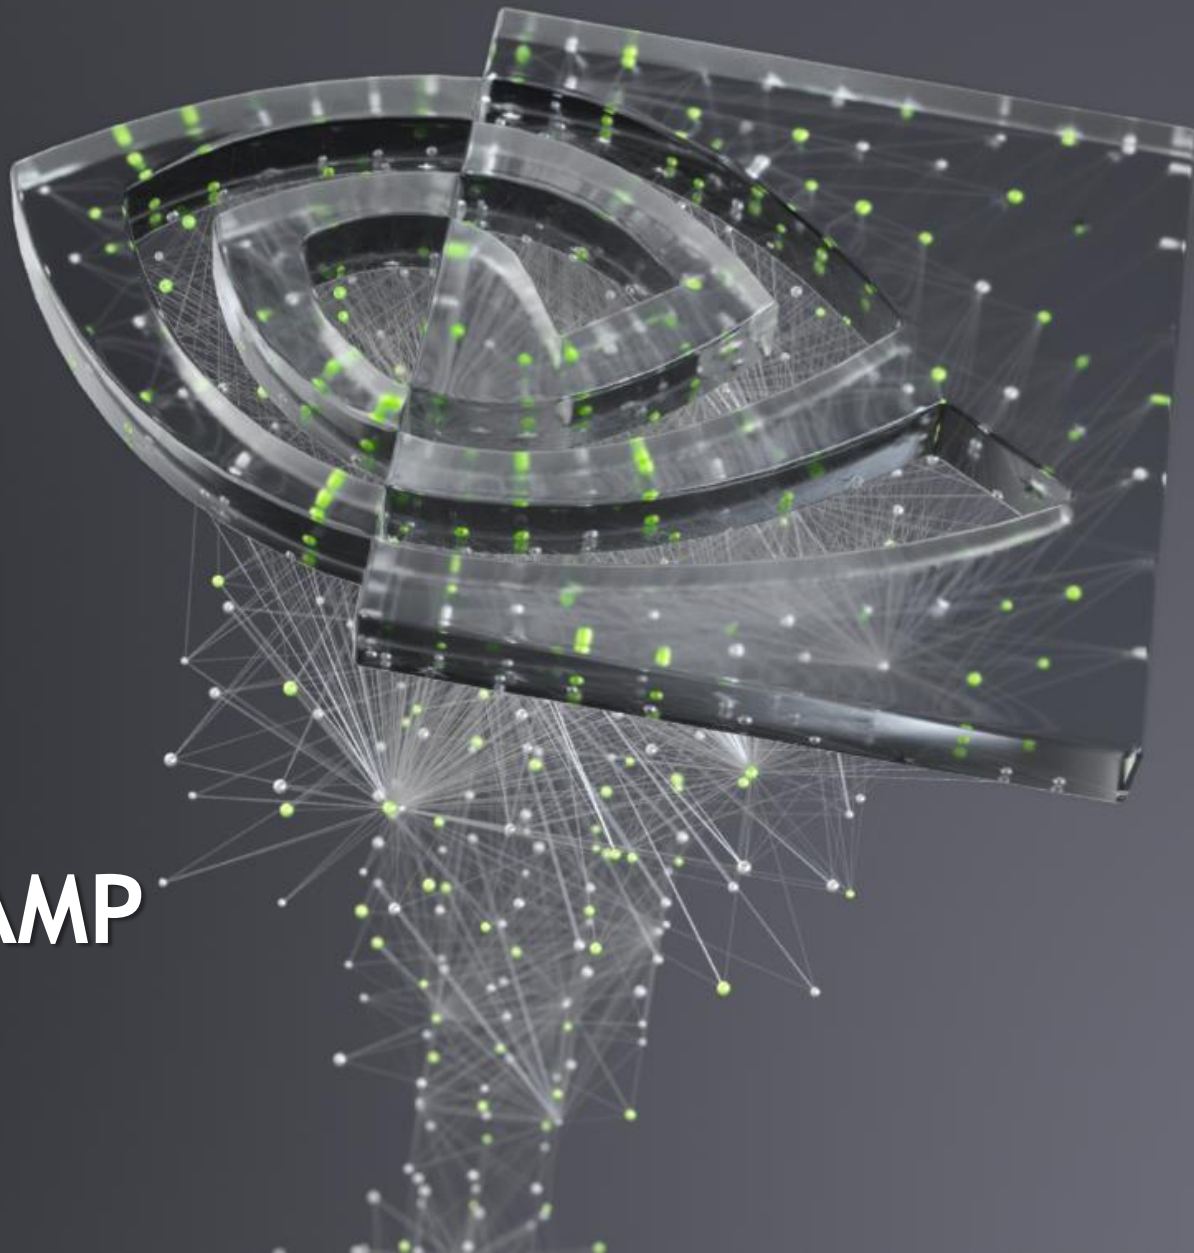




N-WAYS GPU BOOTCAMP

NSIGHT SYSTEM PROFILING



AGENDA

일정표:

Agenda:

1일차 (Intro): 2021년 7월 26일 (월) 9시 - 16시

- 강사 소개 (Moderator) : 09:00 AM - 09:15AM (Lecture)
- 클러스터 연결하기 (okta login) : 09:15 AM - 09:30 AM (Lecture)
- GPU 컴퓨팅 개요 : 09:30 AM - 10:15 AM (Lecture)
- 프로파일링 툴 개요 : 10:15 AM - 11:00 AM (Lecture)
- 점심 시간 : 11:00 AM - 12:30 PM
- 클러스터 연결하기 : 12:30 PM - 13:00 PM (Lecture)
- OpenMP/OpenACC 기반 가속 : 13:00 PM - 15:30 PM (Lecture + Lab)
- Std C++ , Fortran 기반 가속 : 15:30 PM - 16:00 PM (Lecture + Lab)

Day 2 (Advanced): 2021년 7월 27일 (화) 9시 - 15시 30분

- 클러스터 연결하기 (okta login) : 09:00 AM - 09:15 AM (Lecture)
- CUDA C 프로그래밍 : 09:15 AM - 11:00 AM (Lecture + Lab)
- 점심 시간 : 11:00 AM - 12:30 PM
- 미니 챌린지 : 12:30 PM - 14:30 PM (Lab)
- 마무리 : 14:30 PM - 15:30 PM (Feedback)


NSIGHT SYSTEM

What to expect?

- Basic introduction to Nsight family of tools
- Using Nsight Systems

DEVELOPMENT CYCLE

- **Analyze** your code to determine most likely places needing parallelization or optimization.
- **Parallelize** your code by starting with the most time consuming parts and check for correctness.
- **Optimize** your code to improve observed speed-up from parallelization.



PROFILING SEQUENTIAL CODE

Profile Your Code


Obtain detailed information about how the code ran.

This can include information such as:

- Total runtime
- Runtime of individual routines
- Hardware counters

Identify the portions of code that took the longest to run. We want to focus on these “hotspots” when parallelizing.

Hotspot Analysis





NVIDIA NSIGHT FAMILY

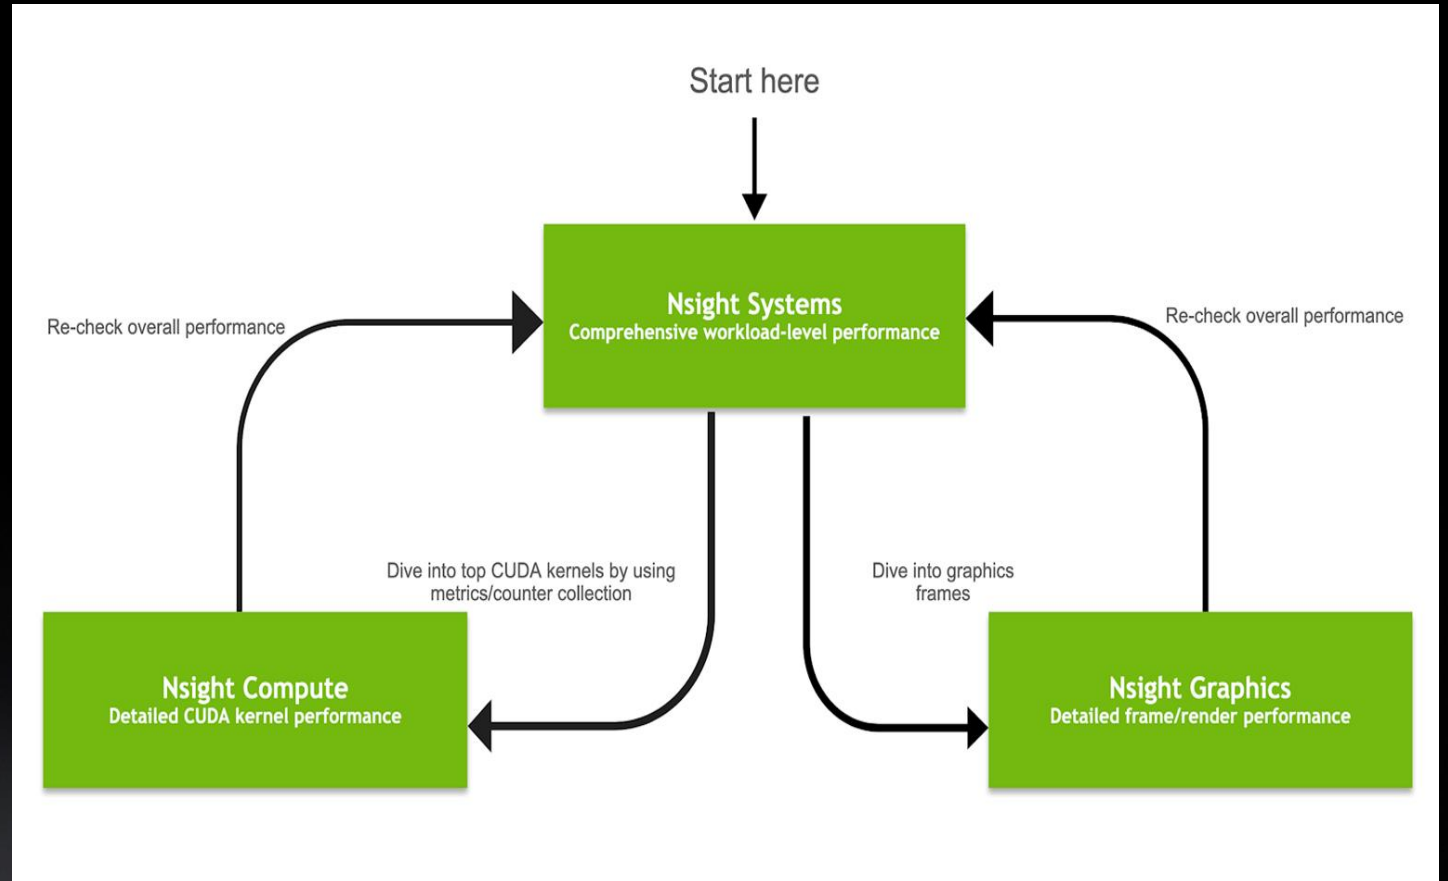
Nsight Product Family

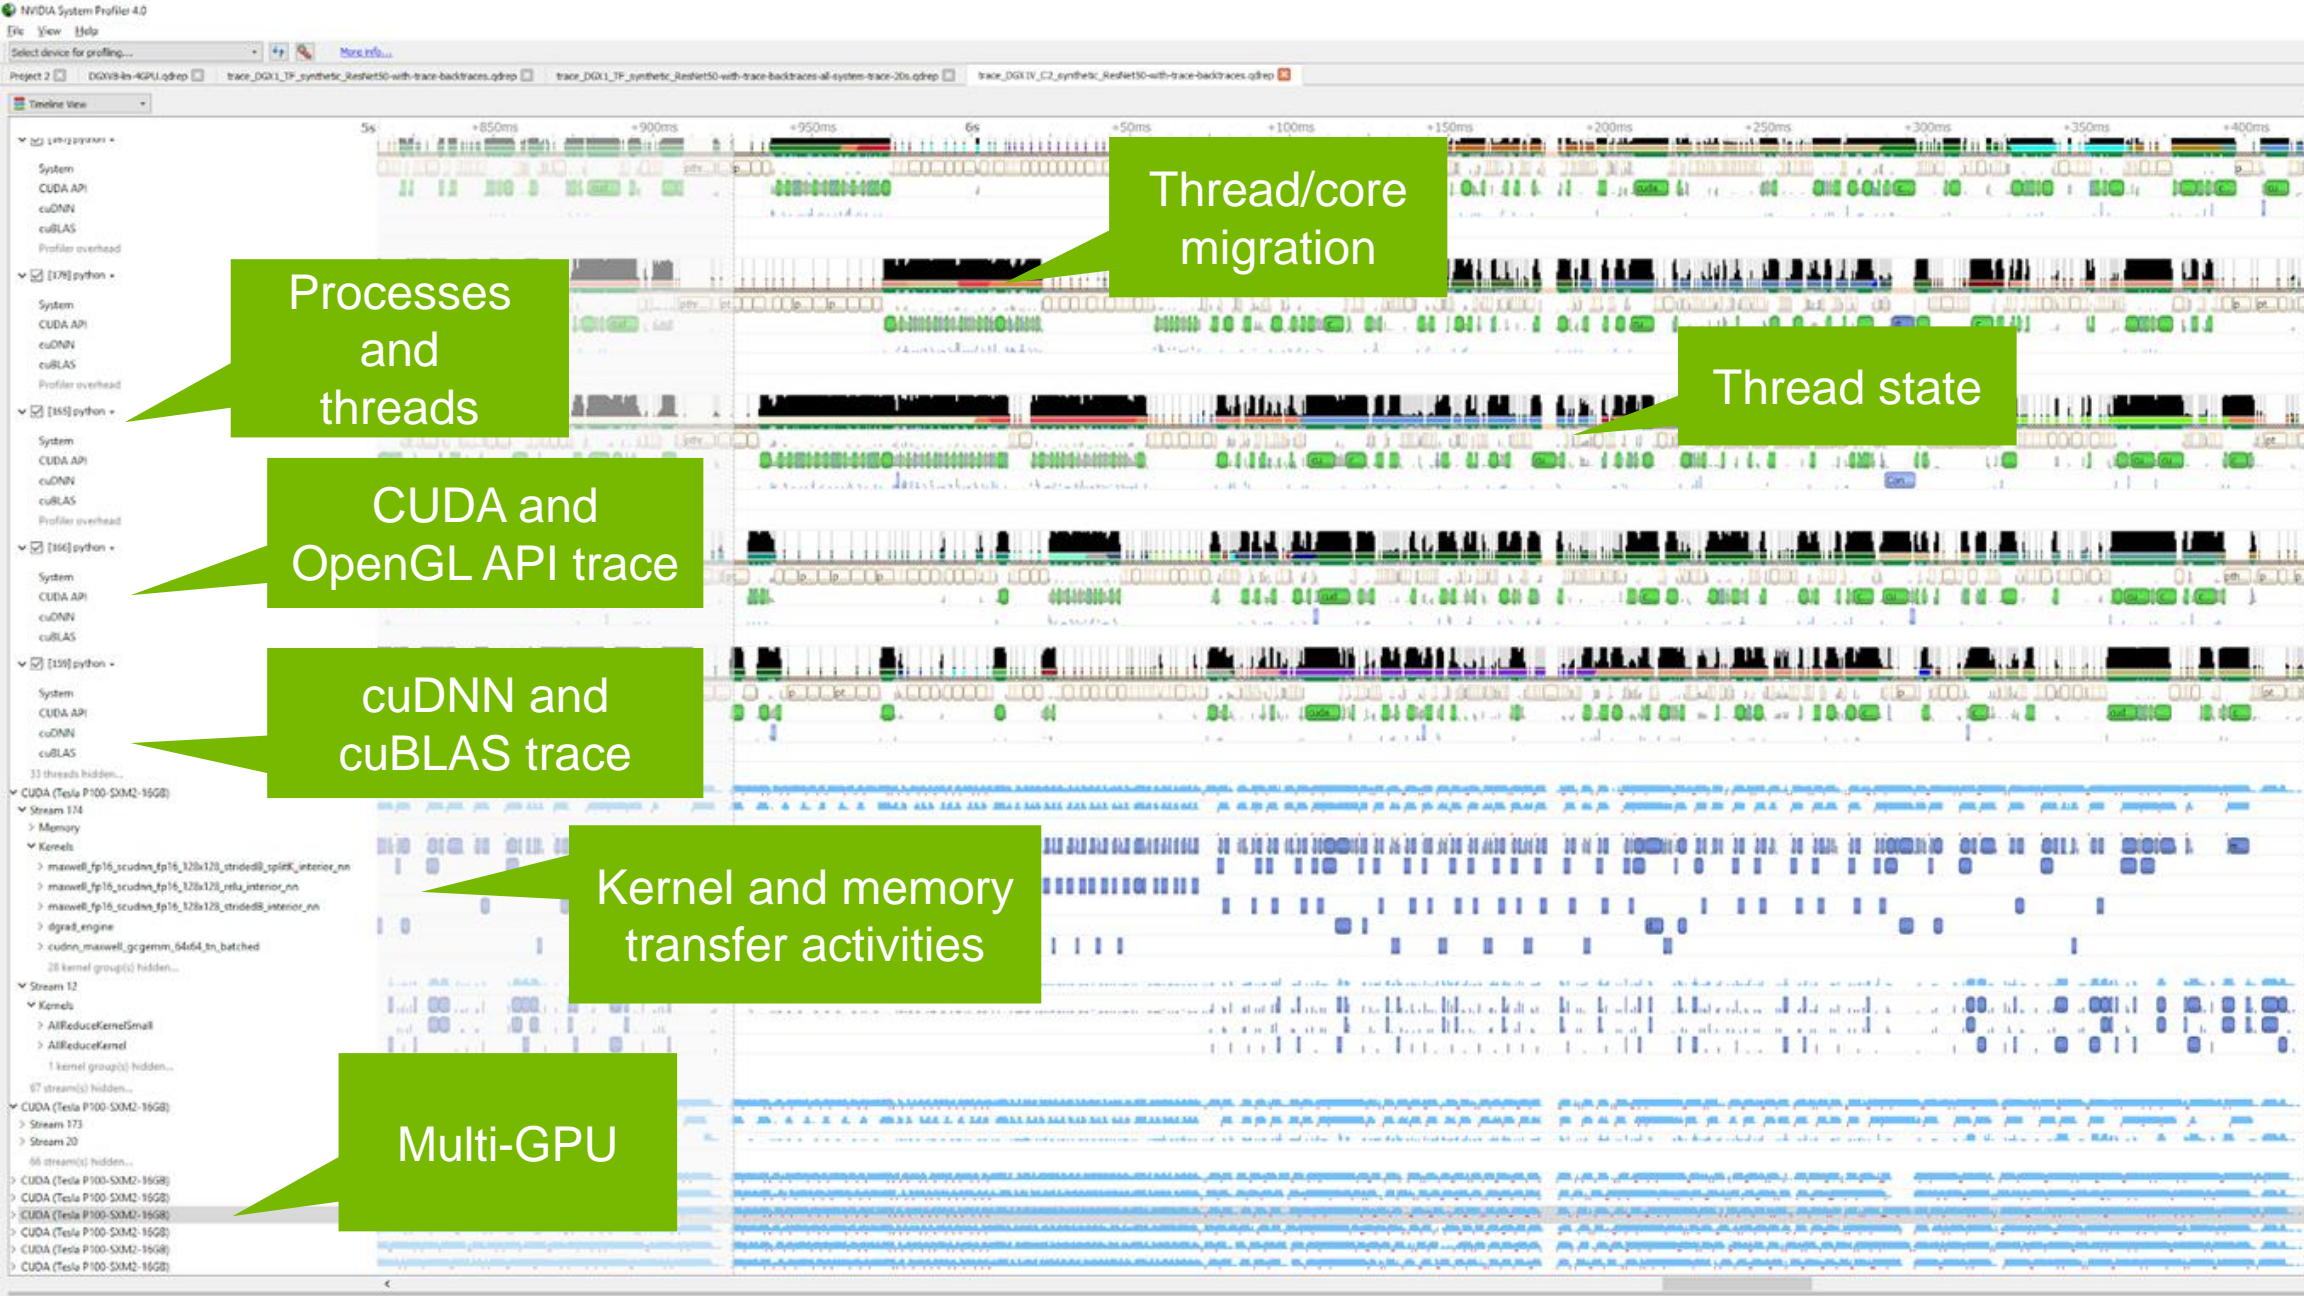
Workflow

Nsight Systems - Analyze application algorithm system-wide

Nsight Compute - Debug/optimize CUDA kernel

Nsight Graphics - Debug/optimize graphics workloads





Thread/cycle migration

Processes and threads

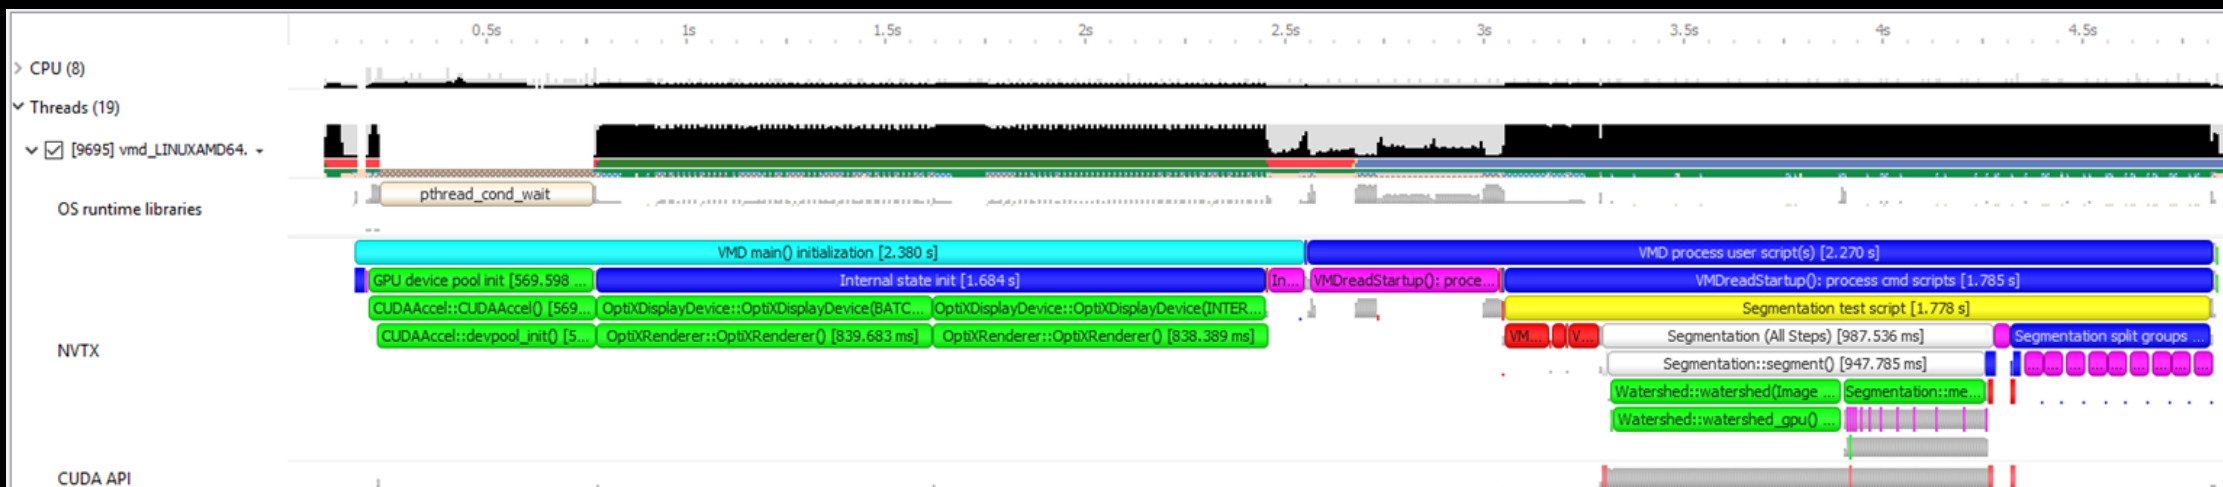
Thread state

CUDA and OpenGL API trace

cuDNN and cuBLAS trace

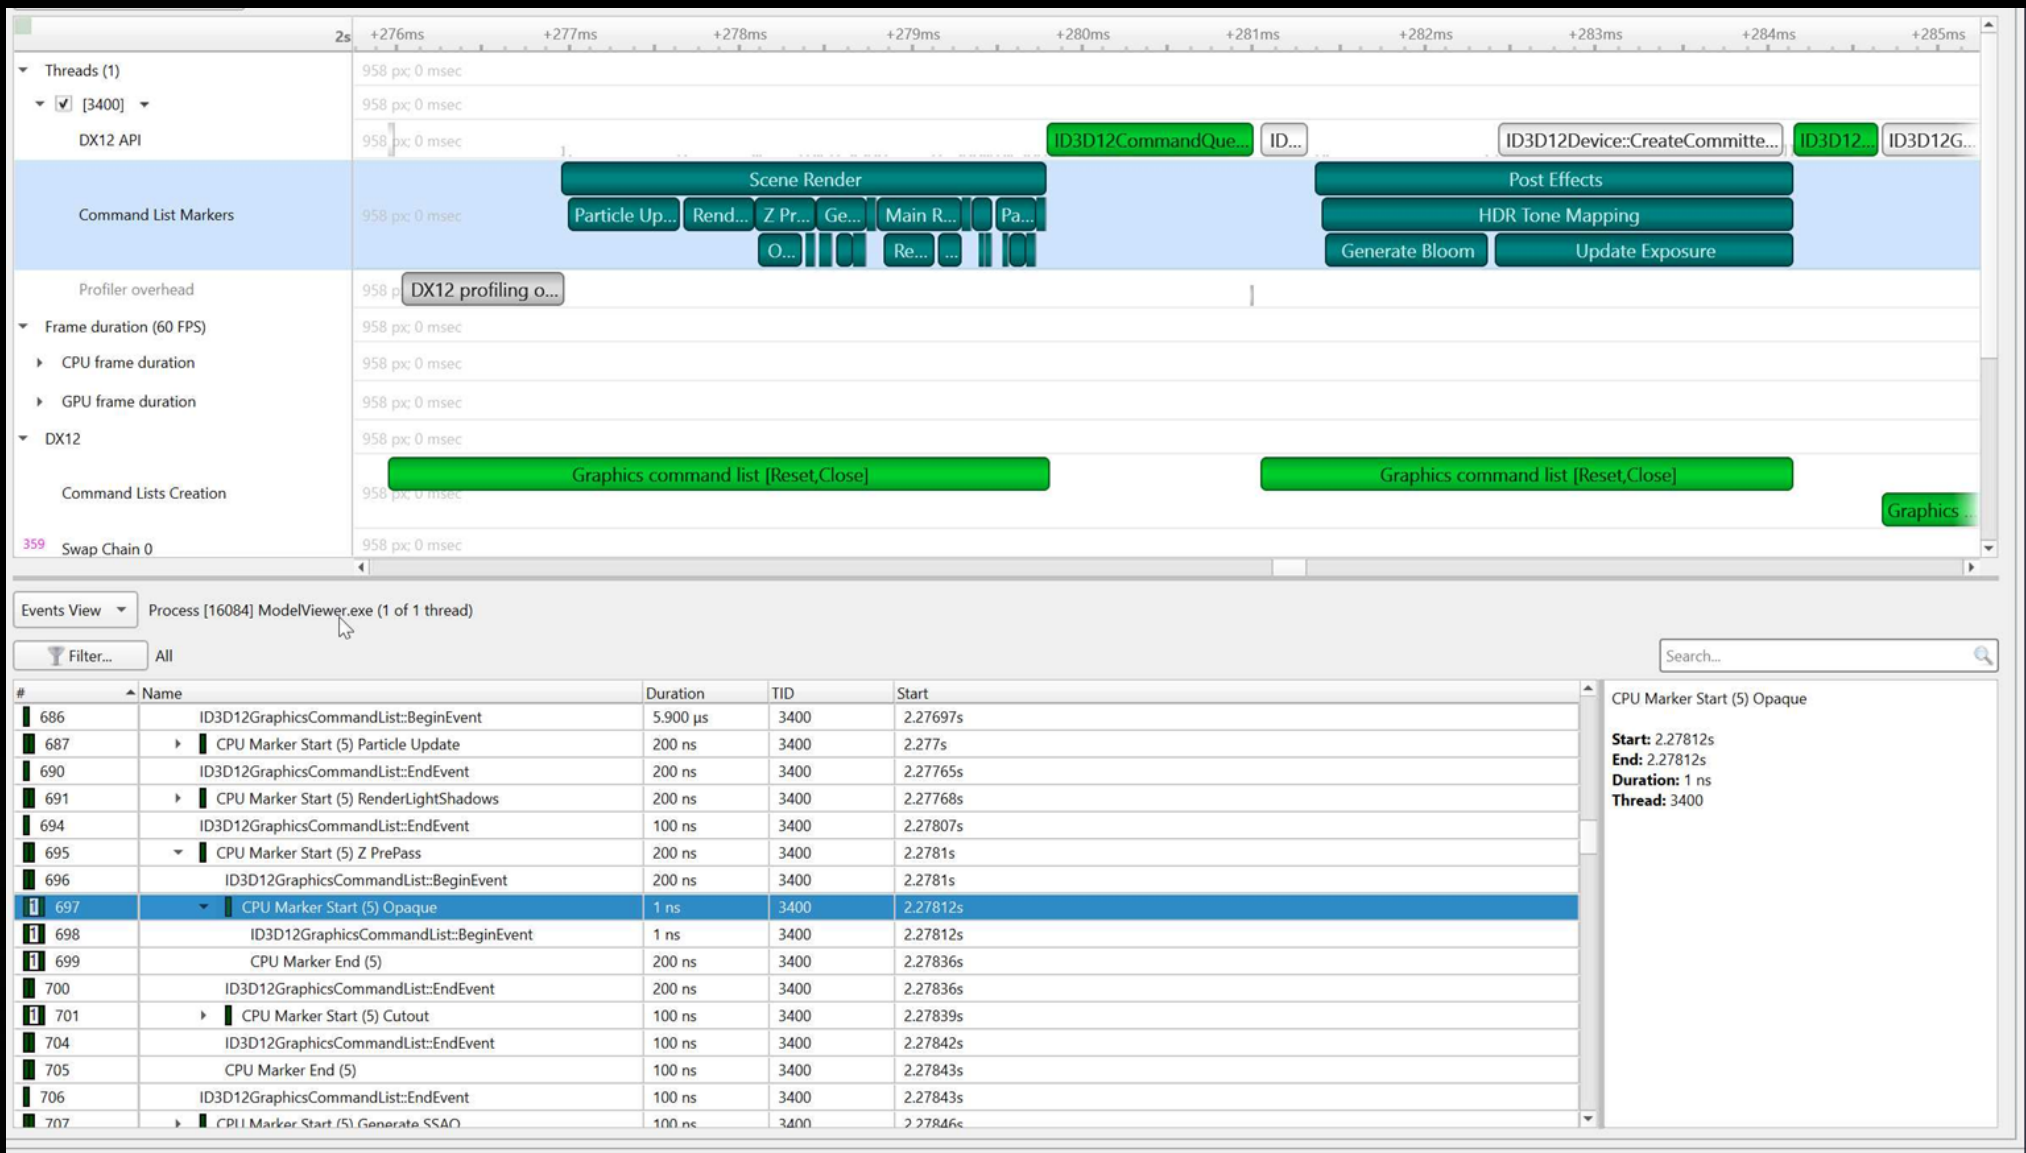
Kernel and memory transfer activities

Multi-GPU



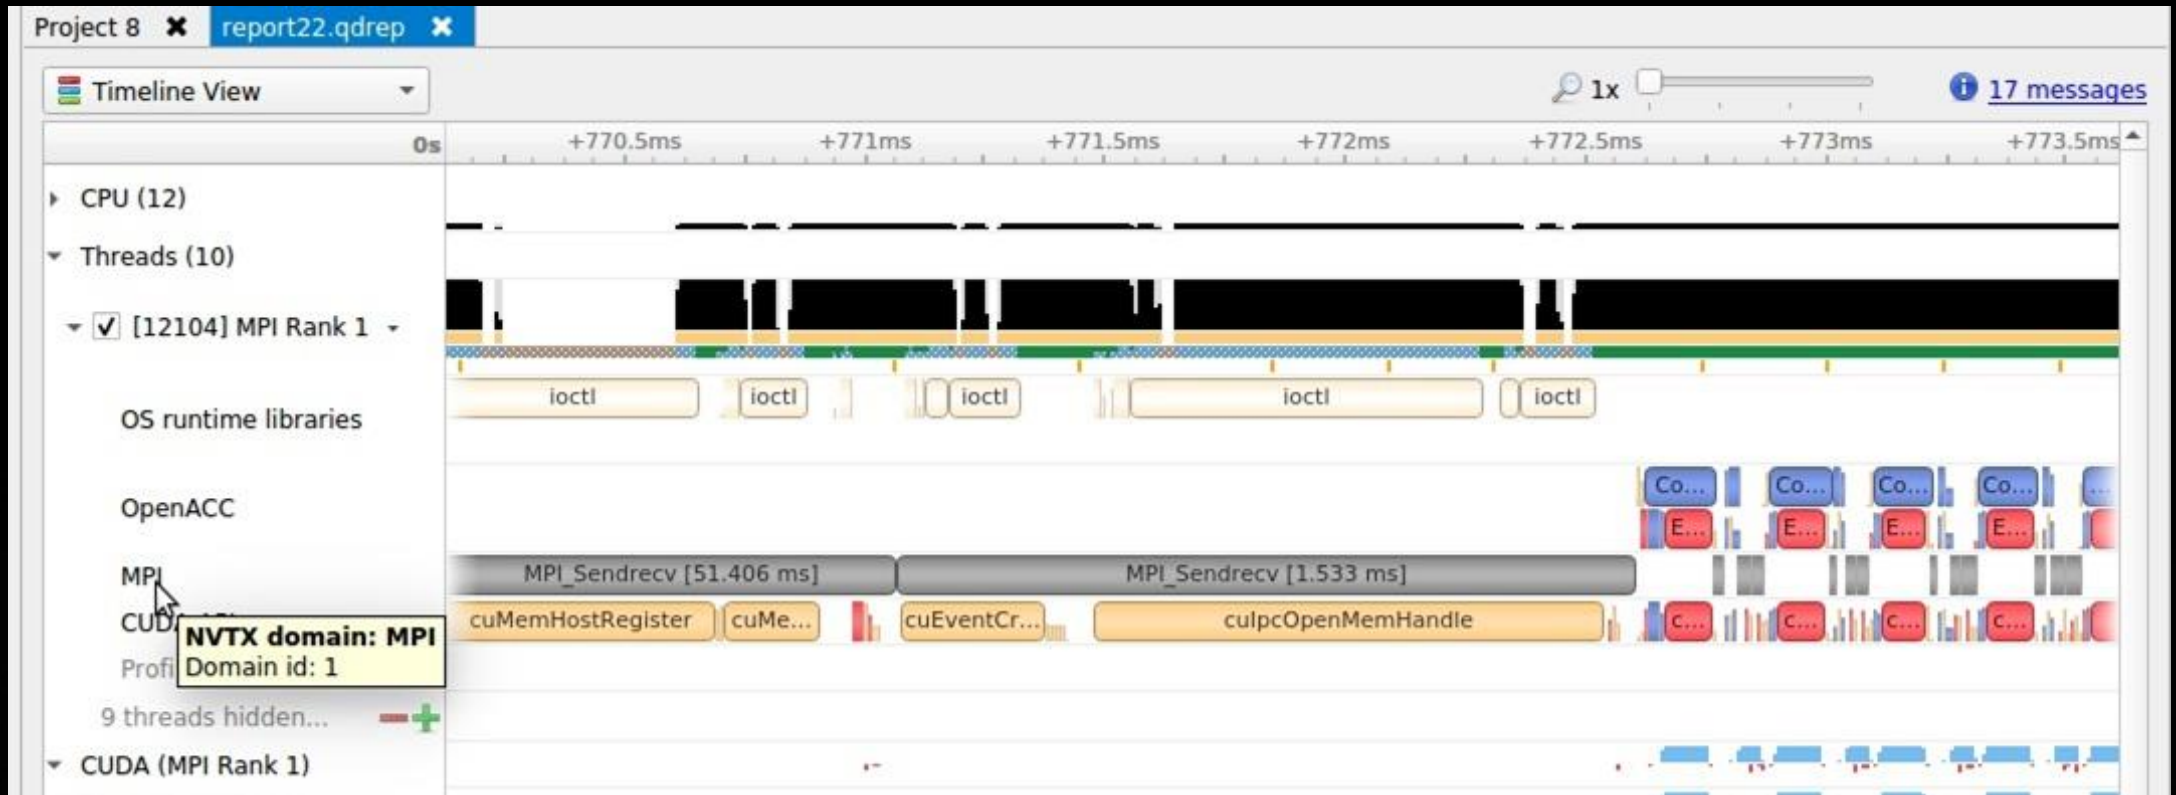
USER ANNOTATIONS APIS FOR CPU & GPU NVTX, OPENGGL, VULKAN, AND DIRECT3D PERFORMANCE MARKERS


EXAMPLE: VISUAL MOLECULAR DYNAMICS (VMD) ALGORITHMS VISUALIZED WITH NVTX ON CPU



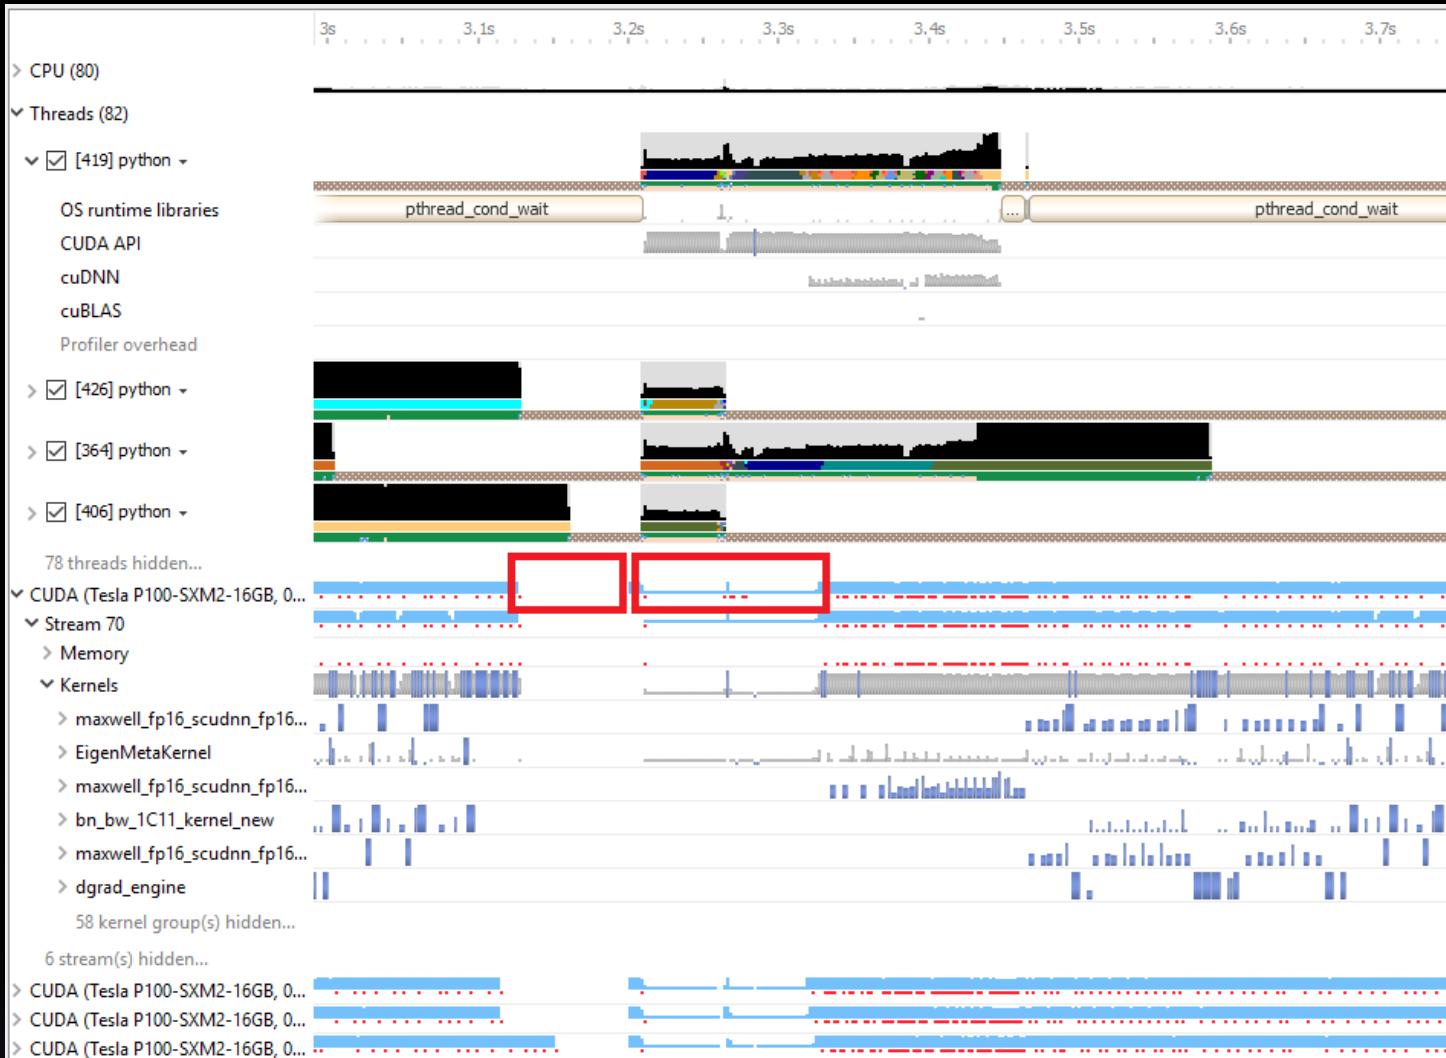
EVENT TABLE

MPI & OPENACC TRACE



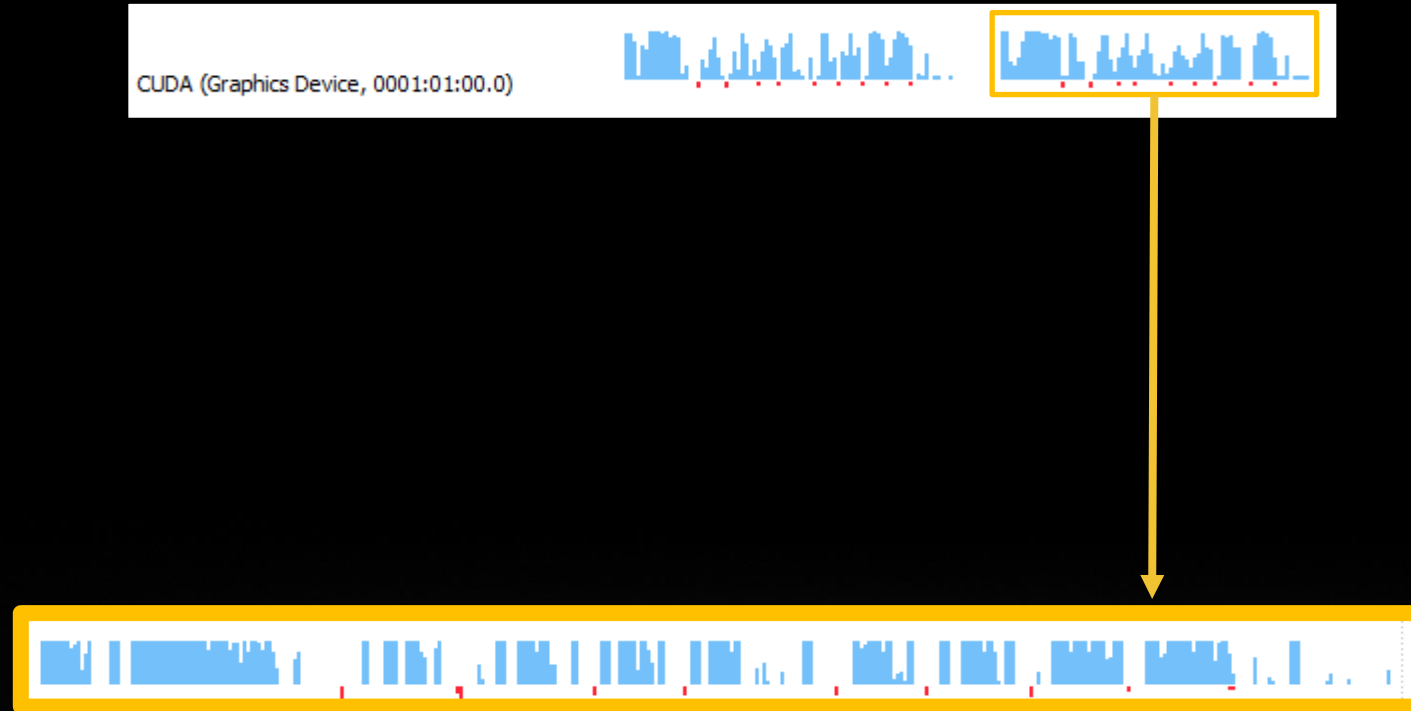


GPU API LAUNCH TO HW WORKLOAD CORRELATION

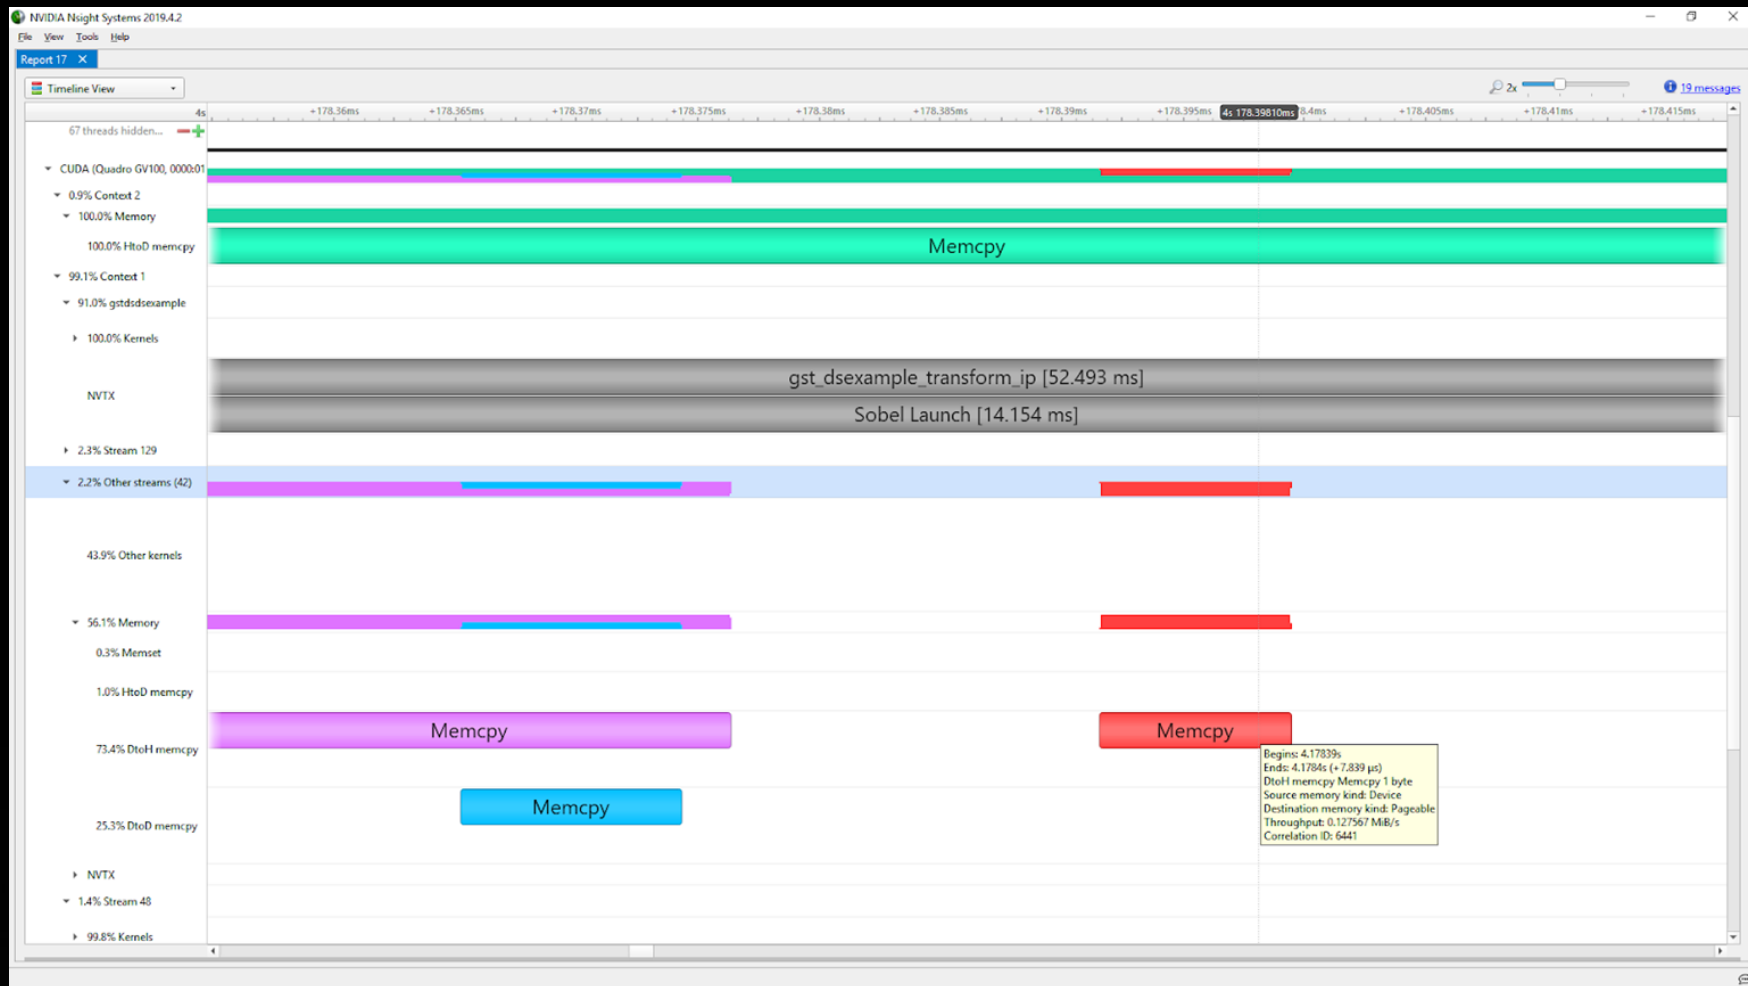


GPU IDLE AND LOW UTILIZATION LEVEL OF DETAIL

GPU UTILIZATION BASED ON PERCENTAGE TIME COVERAGE



ZOOMING IN REVEALS GAPS WHERE THERE WERE VALLEYS



CUDA MEMORY TRANSFER COLOR PALLETTE SHOW DIRECTION AND PAGEABLE MEMORY HAZARDS



CUDA UNIFIED VIRTUAL MEMORY (UVM) TRANSFERS



PROFILING WITH NSIGHT SYSTEM AND NVTX

PROFILING SEQUENTIAL CODE

Using Command Line Interface (CLI)

NVIDIA Nsight Systems CLI provides

- Simple interface to collect data
- Can be copied to any system and analysed later
- Profiles both serial and parallel code
- For more info enter `nsys --help` on the terminal

To profile a serial application with NVIDIA Nsight Systems, we use NVIDIA Tools Extension (NVTX) API functions in addition to collecting backtraces while sampling.

PROFILING SEQUENTIAL CODE

NVIDIA Tools Extension API (NVTX) library

What is it?

- A C-based Application Programming Interface (API) for annotating events
- Can be easily integrated to the application
- Can be used with NVIDIA Nsight Systems

Why?

- Allows manual instrumentation of the application
- Allows additional information for profiling (e.g: tracing of CPU events and time ranges)

How?

- Import the header only C library `nvToolsExt.h`
- Wrap the code region or a specific function with `nvtxRangePush()` and `nvtxRangePop()`

```

#include <string.h>
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include "laplace2d.h"
#include <nvtx3/nvToolsExt.h>

int main(int argc, char** argv)
{
    const int n = 4096;
    const int m = 4096;
    const int iter_max = 1000;

    const double tol = 1.0e-6;
    double error = 1.0;

    double *restrict A = (double*)malloc(sizeof(double)*n*m);
    double *restrict Anew = (double*)malloc(sizeof(double)*n*m);

    nvtxRangePushA("init");
    initialize(A, Anew, m, n);
    nvtxRangePop();

    printf("Jacobi relaxation Calculation: %d x %d mesh\n", n, m);

    double st = omp_get_wtime();
    int iter = 0;

    nvtxRangePushA("while");
    while ( error > tol && iter < iter_max )
    {
        nvtxRangePushA("calc");
        error = calcNext(A, Anew, m, n);
        nvtxRangePop();

        nvtxRangePushA("swap");
        swap(A, Anew, m, n);
        nvtxRangePop();

        if(iter % 100 == 0) printf("%5d, %0.6f\n", iter, error);

        iter++;
    }
    nvtxRangePop();

    double runtime = omp_get_wtime() - st;

    printf(" total: %f s\n", runtime);

    deallocate(A, Anew);

    return 0;
}

```

jacobi.c
(starting and ending of ranges are
highlighted with the same color)

- t** Selects the APIs to be traced (nvtx in this example)
- status** if true, generates summary of statistics after the collection
- b** Selects the backtrace method to use while sampling. The option dwarf uses DWARF's CFI (Call Frame Information).
- force-overwrite** if true, overwrites the existing results
- o** sets the output (qdrep) filename

```

mozhgank@prm-dgx-32:~/Code/openacc-training-materials/labs/module4/English/C/solutions/parallel$ nsys profile -t nvtx --stats=true -b dwarf --force-overwrite true -o laplace-seq ./laplace-seq
Collecting data...
Jacobi relaxation calculation: 4096 x 4096 mesh
0, 0.250000
100, 0.002397
200, 0.001204
300, 0.000804
400, 0.000603
500, 0.000483
600, 0.000403
700, 0.000345
800, 0.000302
900, 0.000269
total: 55.754501 s
Processing events...
Capturing symbol files...
Saving intermediate "/home/mozhgank/Code/openacc-training-materials/labs/module4/English/C/solutions/parallel/laplace-seq.qdstrm" file to disk...

Importing [=====100%]
Saved report file to "/home/mozhgank/Code/openacc-training-materials/labs/module4/English/C/solutions/parallel/laplace-seq.qdrep"
Exporting 70802 events: [=====100%]

Exported successfully to
/home/mozhgank/Code/openacc-training-materials/labs/module4/English/C/solutions/parallel/laplace-seq.sqlite
Generating NVTX Push-Pop Range Statistics...
NVTX Push-Pop Range Statistics (nanoseconds)

Time(%)  Total Time  Instances  Average  Minimum  Maximum  Range
-----
49.9     55754497966      1  55754497966.0  55754497966  55754497966  while
26.5     29577817696    1000  29577817.7    29092956    65008545    calc
23.4     26163892482    1000  26163892.5    25761418    60129514    swap
0.1       137489808      1   137489808.0   137489808    137489808    init

```

NVTX range
statistics

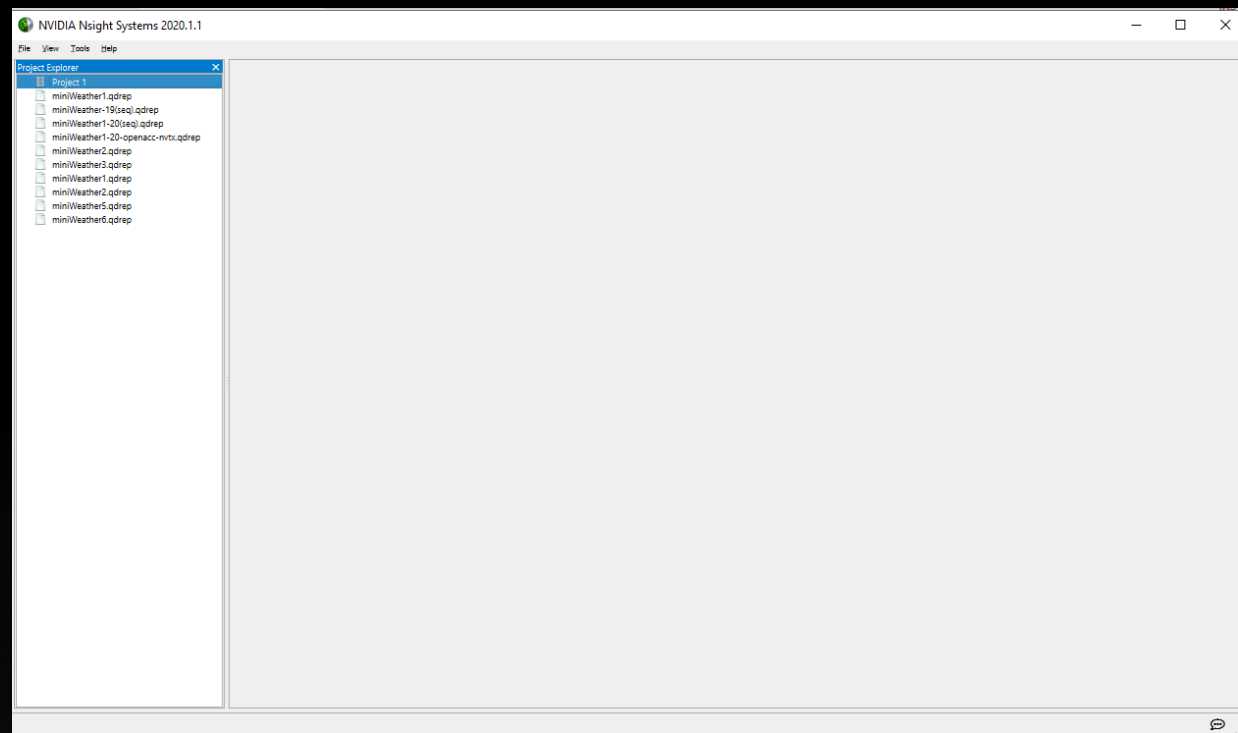
"calc" region (calcNext function) takes 26.6%
"swap" region (swap function) takes 23.4% of
total execution time

Open laplace-seq.qdrep with
Nsight System GUI to view the
timeline

PROFILING CODE

Open the generated report files (*.qdrep) from command line in the Nsight Systems profiler.

File > Open



Using Nsight Systems

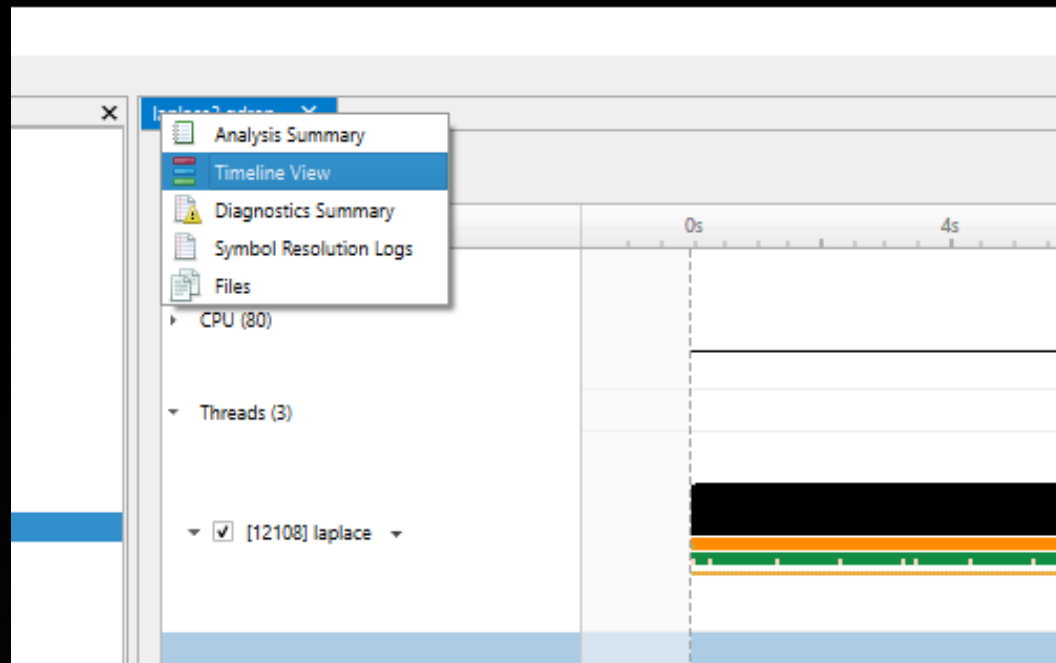
PROFILING CODE

Navigate through the “view selector”.
Using Nsight Systems

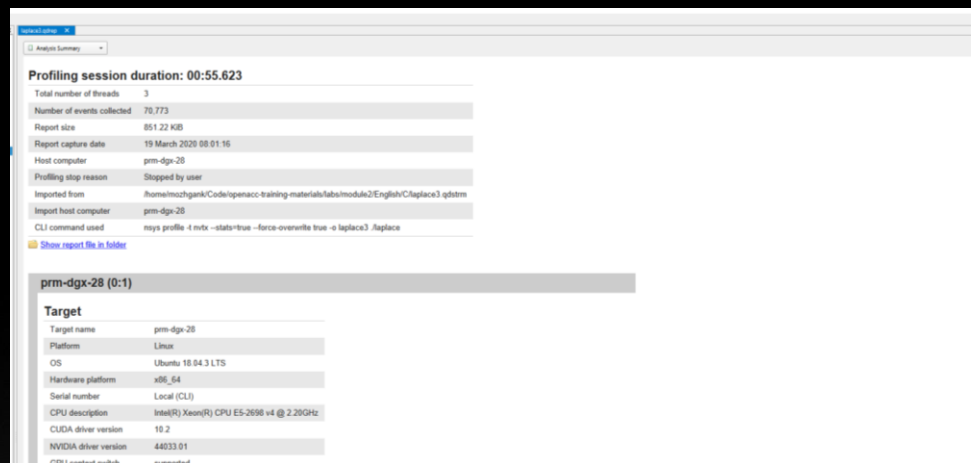
“Analysis summary” shows a summary of the profiling session. To review the project configuration used to generate this report, see next slide.

“Timeline View” contains the timeline at the top, and a bottom pane that contains the events view and the function table.

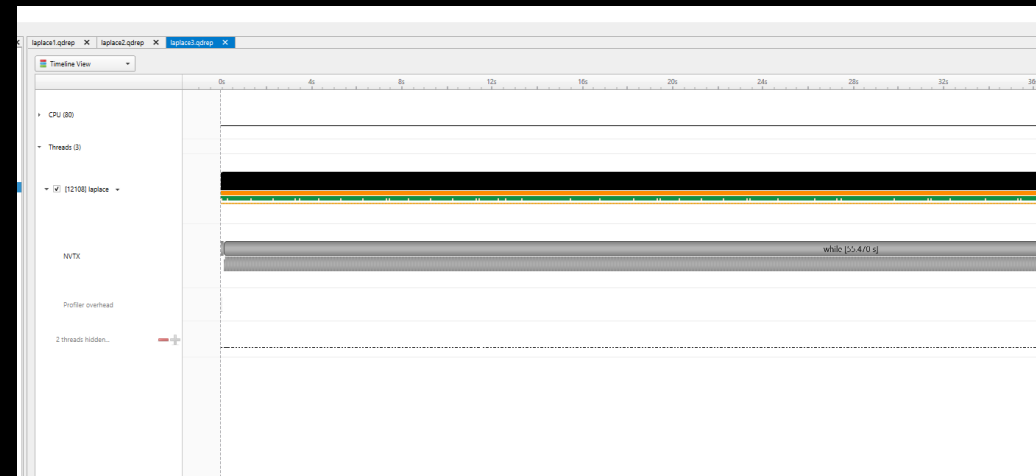
Read more: <https://docs.nvidia.com/nsight-systems>




PROFILING CODE



Analysis Summary



Timeline view
(charts and the hierarchy on the top pane)




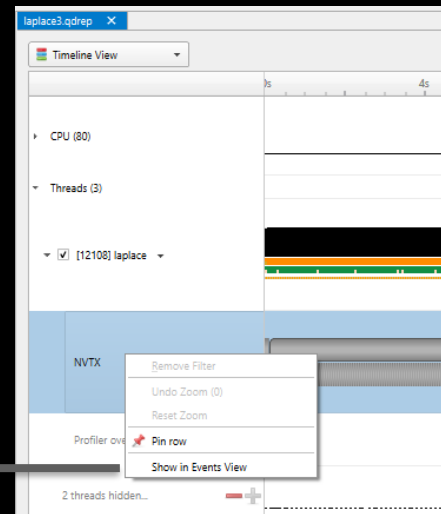
Timeline view
(event view and function table on the bottom pane)

PROFILING CODE

Viewing captured NVTX events and time ranges via Nsight Systems GUI

From the Timeline view, right click on the “NVTX” from the top pane and choose “Show in Events View”.

From the bottom pane, you can now see name of the events captured with the duration.



REFERENCES

<https://docs.nvidia.com/nsight-systems>

<https://developer.nvidia.com/hpc-sdk>



THANK YOU



nvidia.